



UNIVERSITÀ
DEGLI STUDI
DI PADOVA



Streaming processing of cosmic rays using Drift Tubes detectors

Buriola Lorenzo
Campesan Giulia
Giorgetti Sabrina

- Project intro
 - Cluster creation
 - Kafka and Spark settings for the VM
- Kafka and Spark implementations
 - Kafka producer + consumer
 - Kafka/Spark interface
- Analysis of data in analysis
 - Spark to Kafka
- Results

Implemented in a cluster of computers

INPUT - DATA

Data collected by the DAQ of the
Drift Tubes detectors

PROCESSING

Processing of the data performed in
a distributed framework

OUTPUT - RESULTS

Plotting and display live
updates of analysed data

Cluster configuration



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

On Cloud Veneto we had 5 VM, the following steps have been performed in order to configure the cluster:

1. Access VM:

```
$ ssh -L3232:localhost:3232 -L2080:10.67.22.248:8080 -  
L2040:10.67.22.248:4040 user@gate.cloudveneto.it
```

```
$ ssh -L3232:localhost:3232 root@10.67.22.248
```

```
$ jupyter notebook --ip 127.0.0.1 --port 3232 --no-  
browser --allow-root
```

We used the master 10.67.22.248 as our main machine to run Jupyter

SPARK

Master

MAPD-B_Gr12-1 10.67.22.248

Slaves

MAPD-B_Gr12-2 10.67.22.110

MAPD-B_Gr12-3 10.67.22.177

MAPD-B_Gr12-5 10.67.22.16

KAFKA BROKER + ZOOKEEPER

MAPD-B_Gr12-4 10.67.22.185

Cluster configuration



On Cloud Veneto we had 5 VM, the following steps have been performed in order to configure the cluster:

2. *\$ yum install openssh-server openssh-client*

3. Setting the machines names:

```
$ sudo vim /etc/hosts
10.67.22.248 master
10.67.22.185 kafka
10.67.22.110 slave01
10.67.22.177 slave02
10.67.22.16  slave03
```

4. All machines were connected with ssh passwordless:

- A. Generate key : *\$ ssh-keygen -t rsa -P ""*
- B. Make key an authorized one cat: *\$ ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys*
- C. Copy key in all machines: *\$ ssh-copy-id root@<IP>*

SPARK

Master

MAPD-B_Gr12-1 10.67.22.248

Slaves

MAPD-B_Gr12-2 10.67.22.110

MAPD-B_Gr12-3 10.67.22.177

MAPD-B_Gr12-5 10.67.22.16

KAFKA BROKER + ZOOKEEPER

MAPD-B_Gr12-4 10.67.22.185

SPARK

1. Install java

```
$ yum install java-1.8.0-openjdk
```

2. Install pyspark and pyarrow

3. Download and install Spark

```
$ mv spark-3.1.1-bin-hadoop3.2 /usr/local/spark
```

```
$ vim ~/.bashrc modify and $ source ~/.bashrc
```

```
export PATH = PATH:/usr/local/spark/bin
```

4. Spark Master Configuration

```
$ cd /usr/local/spark/conf
```

```
$ cp spark-env.sh.template spark-env.sh
```

```
$ vim spark-env.sh
```

```
export SPARK_MASTER_HOST='<MASTER-IP>'
```

```
export JAVA_HOME=<JAVA-PATH>
```

5. Workers Configuration

```
$ vim slaves
```

```
master
```

```
slave01
```

```
slave02
```

```
slave03
```

KAFKA

1. Download and install Kafka

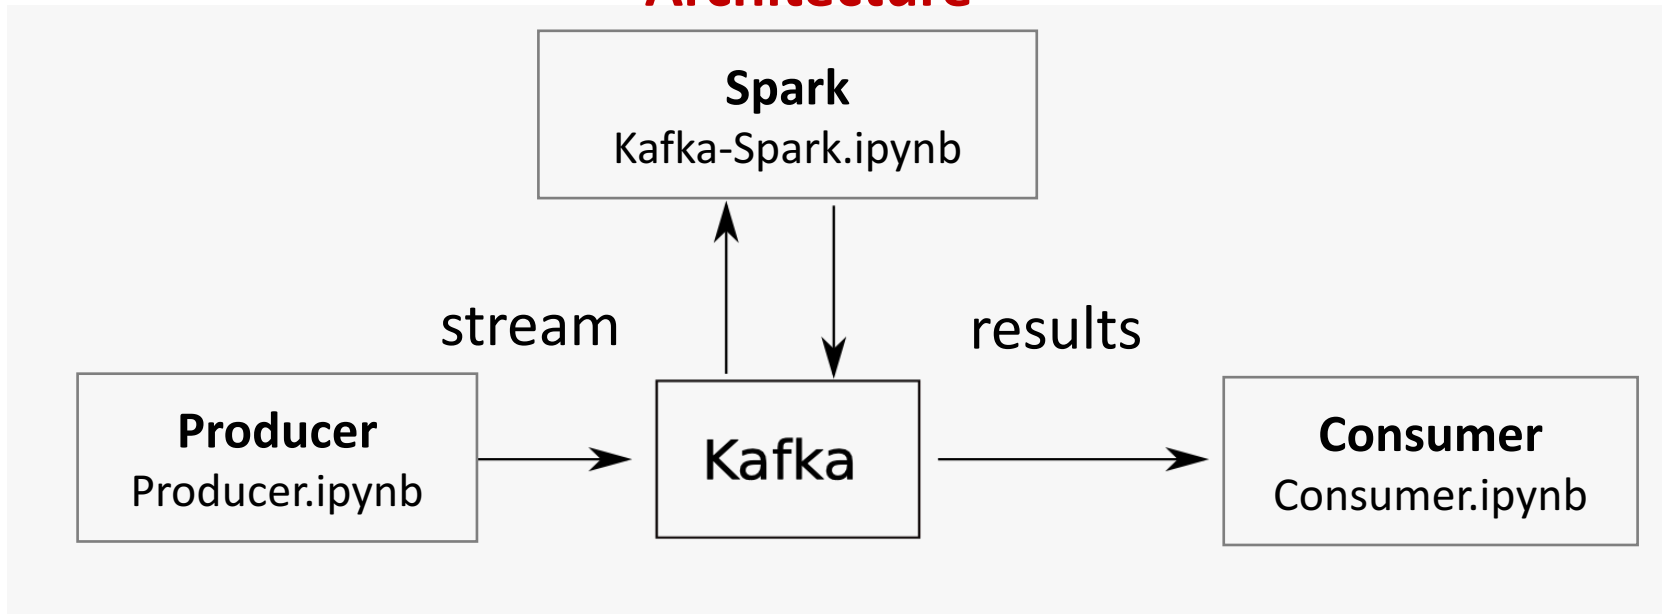
```
kafka_2.13-2.7.0
```

2. Configuration

```
$ vim config/zookeeper.properties
```

```
$ vim config/server.properties
```

Architecture



Starting Kafka from terminal

```
ssh kafka
kafka_2.13-2.7.0/bin/zookeeper-server-start.sh kafka_2.13-2.7.0/config/zookeeper.properties &
kafka_2.13-2.7.0/bin/kafka-server-start.sh kafka_2.13-2.7.0/config/server.properties &
exit
```

Producer

```
from kafka.admin import KafkaAdminClient
from kafka import KafkaProducer
KAFKA_BOOTSTRAP_SERVERS = ['kafka:9092']
producer = KafkaProducer(bootstrap_servers=KAFKA_BOOTSTRAP_SERVERS)
kafka_admin = KafkaAdminClient(bootstrap_servers=KAFKA_BOOTSTRAP_SERVERS)
```

```
kafka_admin.list_topics()
```

```
['stream', 'results', '__consumer_offsets']
```

Import and sending data

```
ssl._create_default_https_context = ssl._create_unverified_context
```

```
link="https://cloud-areapd.pd.infn.it:5210/swift/v1/AUTH_d2e941ce4b324467b6b3d467a923a9bc/MAPD_miniDT_stream/"
```

```
filenames = pd.read_csv(link, delimiter=" ", header=None)
```

```
for i in range(1):
    url= "https://cloud-areapd.pd.infn.it:5210/swift/v1/AUTH_d2e941ce4b324467b6b3d467a923a9bc/MAPD_miniDT_stream/" + str(filenames.loc[i,0])
    k=pd.read_csv(url)
    for j in range(len(k)) :
        df_json = json.loads(k.loc[[j]].to_json(orient='records'))[0]
        producer.send('stream', json.dumps(df_json).encode('utf-8'))
        producer.flush()
        time.sleep(0.001)
```


Kafka topics

```
from kafka import KafkaProducer
from kafka.admin import KafkaAdminClient, NewTopic
```

```
KAFKA_BOOTSTRAP_SERVERS = 'kafka:9092'
```

```
producer = KafkaProducer(bootstrap_servers=KAFKA_BOOTSTRAP_SERVERS)
```

```
kafka_admin = KafkaAdminClient(bootstrap_servers=KAFKA_BOOTSTRAP_SERVERS)
```

```
kafka_admin.delete_topics(['stream', 'results'])
```

```
kafka_admin.list_topics()
```

```
['stream', '__consumer_offsets']
```

```
stream_topic = NewTopic(name='stream',
                        num_partitions=1,
                        replication_factor=1)
```

```
results_topic = NewTopic(name='results',
                        num_partitions=1,
                        replication_factor=1)
```

```
kafka_admin.create_topics(new_topics=[stream_topic, results_topic])
```

```
kafka_admin.list_topics()
```

```
['stream', '__consumer_offsets']
```

Starting Spark session

```
import findspark
findspark.init('/usr/local/spark/')
```

```
! $SPARK_HOME/sbin/start-all.sh
```

```
starting org.apache.spark.deploy.master.Master, logging to /usr/local/spark/logs/spark-root-org.apache.spark
.deploy.master.Master-1-mapd-b-gr12-1.novalocal.out
master: starting org.apache.spark.deploy.worker.Worker, logging to /usr/local/spark/logs/spark-root-org.apac
e.spark.deploy.worker.Worker-1-mapd-b-gr12-1.novalocal.out
slave02: starting org.apache.spark.deploy.worker.Worker, logging to /usr/local/spark/logs/spark-root-org.apac
he.spark.deploy.worker.Worker-1-mapd-b-gr12-3.novalocal.out
slave03: starting org.apache.spark.deploy.worker.Worker, logging to /usr/local/spark/logs/spark-root-org.apac
he.spark.deploy.worker.Worker-1-mapd-b-gr12-5.novalocal.out
slave01: starting org.apache.spark.deploy.worker.Worker, logging to /usr/local/spark/logs/spark-root-org.apac
he.spark.deploy.worker.Worker-1-mapd-b-gr12-2.novalocal.out
```

```
from pyspark.sql import SparkSession
```

```
spark = SparkSession.builder \
    .master("spark://master:7077") \
    .appName("Project_MAPDB_application") \
    .config("spark.sql.execution.arrow.pyspark.enabled", "true") \
    .config("spark.sql.execution.arrow.pyspark.fallback.enabled", "false") \
    .config("spark.jars.packages", "org.apache.spark:spark-sql-kafka-0-10_2.12:3.1.1") \
    .getOrCreate()
```

```
sc = spark.sparkContext
sc
```

SparkContext

[Spark UI](#)

Version

v3.1.1

Master

spark://master:7077

AppName

Project_MAPDB_application

Kafka - Spark



```
inputDF = spark.readStream\  
    .format("kafka")\  
    .option("kafka.bootstrap.servers", KAFKA_BOOTSTRAP_SERVERS)\  
    .option('subscribe', 'stream')\  
    .load()
```

```
from pyspark.sql.functions import from_json, col  
from pyspark.sql.types import StructField, StructType, StringType, DoubleType, IntegerType, LongType  
  
## the schema of the json data format used to create the messages  
schema = StructType(  
    [  
        StructField("HEAD", IntegerType()),  
        StructField("FPGA", StringType()),  
        StructField("TDC_CHANNEL", IntegerType()),  
        StructField("ORBIT_CNT", StringType()),  
        StructField("BX_COUNTER", StringType()),  
        StructField("TDC_MEAS", StringType())  
    ]  
)  
## a new DF can be created from the previous by using the pyspark.sql functions  
jsonDF = inputDF.select(from_json(col("value").alias('value')).cast("string"), schema).alias('value'))
```

```
flatDF = jsonDF.selectExpr(  
    "value.HEAD",  
    "value.FPGA",  
    "value.TDC_CHANNEL",  
    "value.ORBIT_CNT",  
    "value.BX_COUNTER",  
    "value.TDC_MEAS")
```

```
inputDF.printSchema()
```

```
root  
|-- key: binary (nullable = true)  
|-- value: binary (nullable = true)  
|-- topic: string (nullable = true)  
|-- partition: integer (nullable = true)  
|-- offset: long (nullable = true)  
|-- timestamp: timestamp (nullable = true)  
|-- timestampType: integer (nullable = true)
```

```
jsonDF.printSchema()
```

```
root  
|-- value: struct (nullable = true)  
|   |-- HEAD: integer (nullable = true)  
|   |-- FPGA: string (nullable = true)  
|   |-- TDC_CHANNEL: integer (nullable = true)  
|   |-- ORBIT_CNT: string (nullable = true)  
|   |-- BX_COUNTER: string (nullable = true)  
|   |-- TDC_MEAS: string (nullable = true)
```

```
flatDF.printSchema()
```

```
root  
|-- HEAD: integer (nullable = true)  
|-- FPGA: string (nullable = true)  
|-- TDC_CHANNEL: integer (nullable = true)  
|-- ORBIT_CNT: string (nullable = true)  
|-- BX_COUNTER: string (nullable = true)  
|-- TDC_MEAS: string (nullable = true)
```



```
def analysis(df, epoch_id):  
  
    #total events  
    tot = df.count()  
  
    #clean  
    df_clean = df.where(col('HEAD')==2)  
  
    #point 1  
    tot_hits = df_clean.count()  
  
    #chambers  
    chamber_0 = df_clean \  
        .where(col('FPGA') == 0) \  
        .where(col('TDC_CHANNEL') >= 0) \  
        .where(col('TDC_CHANNEL') <= 63)  
  
    chamber_1 = df_clean \  
        .where(col('FPGA') == 0) \  
        .where(col('TDC_CHANNEL') >= 64) \  
        .where(col('TDC_CHANNEL') <= 127)  
  
    chamber_2 = df_clean \  
        .where(col('FPGA') == 1) \  
        .where(col('TDC_CHANNEL') >= 0) \  
        .where(col('TDC_CHANNEL') <= 63) \  
  
    chamber_3 = df_clean \  
        .where(col('FPGA') == 1) \  
        .where(col('TDC_CHANNEL') >= 64) \  
        .where(col('TDC_CHANNEL') <= 127) \  
  
    #point2  
    tot_hits_ch0 = chamber_0.count()  
    tot_hits_ch1 = chamber_1.count()  
    tot_hits_ch2 = chamber_2.count()  
    tot_hits_ch3 = chamber_3.count()
```

Data processing

- Filtering with condition HEAD==2
- Compute total number of processed hits, post-cleansing
- Map into chambers
- Compute total number of processed hits, post-cleansing, per chamber

	HEAD	FPGA	TDC_CHANNEL	ORBIT_CNT	BX_COUNTER	TDC_MEAS
0	1	1	0	3387315431	0	130.000000
1	0	1	2	3387315431	1119	24.000000
2	4	1	0	3387315431	0	-0.573730
3	5	1	0	3387315431	0	45.500000
4	2	0	75	3387200947	2922	2.000000
5	2	0	105	3387200955	2227	29.000000
6	2	0	107	3387200955	2234	7.000000
7	2	0	126	3387200973	476	29.000000
8	2	1	55	3387200955	1797	12.000000
9	2	1	36	3387200956	2165	28.000000
10	2	1	51	3387200970	249	14.000000



#point 3: histogram of the counts of active TDC_CHANNEL

```
ch0_hist_bins = chamber_0.groupBy('TDC_CHANNEL').count().select(col('TDC_CHANNEL')).collect()
ch1_hist_bins = chamber_1.groupBy('TDC_CHANNEL').count().select(col('TDC_CHANNEL')).collect()
ch2_hist_bins = chamber_2.groupBy('TDC_CHANNEL').count().select(col('TDC_CHANNEL')).collect()
ch3_hist_bins = chamber_3.groupBy('TDC_CHANNEL').count().select(col('TDC_CHANNEL')).collect()
```

```
ch0_hist_counts = chamber_0.groupBy('TDC_CHANNEL').count().select(col('count')).collect()
ch1_hist_counts = chamber_1.groupBy('TDC_CHANNEL').count().select(col('count')).collect()
ch2_hist_counts = chamber_2.groupBy('TDC_CHANNEL').count().select(col('count')).collect()
ch3_hist_counts = chamber_3.groupBy('TDC_CHANNEL').count().select(col('count')).collect()
```

#point 4: histogram of the total number of active TDC_CHANNEL in each ORBIT_CNT

```
ch0_hist_orbs_bins=chamber_0.groupBy('ORBIT_CNT').agg(F.countDistinct("TDC_CHANNEL"))\
    .groupBy(col('count(TDC_CHANNEL)').count()).select(col('count(TDC_CHANNEL)').collect()
ch1_hist_orbs_bins=chamber_1.groupBy('ORBIT_CNT').agg(F.countDistinct("TDC_CHANNEL"))\
    .groupBy(col('count(TDC_CHANNEL)').count()).select(col('count(TDC_CHANNEL)').collect()
ch2_hist_orbs_bins=chamber_2.groupBy('ORBIT_CNT').agg(F.countDistinct("TDC_CHANNEL"))\
    .groupBy(col('count(TDC_CHANNEL)').count()).select(col('count(TDC_CHANNEL)').collect()
ch3_hist_orbs_bins=chamber_3.groupBy('ORBIT_CNT').agg(F.countDistinct("TDC_CHANNEL"))\
    .groupBy(col('count(TDC_CHANNEL)').count()).select(col('count(TDC_CHANNEL)').collect()
```

```
ch0_hist_orbs_counts=chamber_0.groupBy('ORBIT_CNT').agg(F.countDistinct("TDC_CHANNEL"))\
    .groupBy(col('count(TDC_CHANNEL)').count()).select(col('count')).collect()
ch1_hist_orbs_counts=chamber_1.groupBy('ORBIT_CNT').agg(F.countDistinct("TDC_CHANNEL"))\
    .groupBy(col('count(TDC_CHANNEL)').count()).select(col('count')).collect()
ch2_hist_orbs_counts=chamber_2.groupBy('ORBIT_CNT').agg(F.countDistinct("TDC_CHANNEL"))\
    .groupBy(col('count(TDC_CHANNEL)').count()).select(col('count')).collect()
ch3_hist_orbs_counts=chamber_3.groupBy('ORBIT_CNT').agg(F.countDistinct("TDC_CHANNEL"))\
    .groupBy(col('count(TDC_CHANNEL)').count()).select(col('count')).collect()
```

Data processing

- Histogram of the counts of active TDC_CHANNEL per chamber
- Histogram of the total number of active TDC_CHANNEL per chamber in each ORBIT_CNT

Streaming of results



```
outputJson = {'tot_import':tot,
              'hits': tot_hits,
              'hitsPerChamber': [tot_hits_ch0, tot_hits_ch1, tot_hits_ch2, tot_hits_ch3],
              'histo_ch0': [ch0_hist_bins, ch0_hist_counts],
              'histo_ch1': [ch1_hist_bins, ch1_hist_counts],
              'histo_ch2': [ch2_hist_bins, ch2_hist_counts],
              'histo_ch3': [ch3_hist_bins, ch3_hist_counts],
              'histo_orbit_ch0': [ch0_hist_orbs_bins, ch0_hist_orbs_counts],
              'histo_orbit_ch1': [ch1_hist_orbs_bins, ch1_hist_orbs_counts],
              'histo_orbit_ch2': [ch2_hist_orbs_bins, ch2_hist_orbs_counts],
              'histo_orbit_ch3': [ch3_hist_orbs_bins, ch3_hist_orbs_counts]
            }

producer.send('results', json.dumps(outputJson).encode('utf-8'))
producer.flush()
pass
```

```
flatDF.isStreaming
```

True

```
flatDF.writeStream\
    .foreachBatch(analysis)\
    .trigger(processingTime='5 seconds')\
    .start()\
    .awaitTermination()
```



Kafka consumer

```
from kafka import KafkaConsumer
KAFKA_BOOTSTRAP_SERVERS = ['kafka:9092']
consumer = KafkaConsumer(bootstrap_servers=KAFKA_BOOTSTRAP_SERVERS, consumer_timeout_ms=20000)
consumer.topics()
```

```
{'results', 'stream'}
```

```
consumer.subscribe('results')
consumer.subscription()
```

```
{'results'}
```

Reading results and plots

```
for message in consumer:
```

```
    mx = json.loads(message.value)
```

```
    print('Number of hits: \t', mx['hits'],
          '\n Number of hits post-cleansing in each chamber: \t %s' %
          str(mx['hitsPerChamber'])
    )
```


```
    name_list = [mx['histo_ch0'], mx['histo_ch1'], mx['histo_ch2'],
                 mx['histo_ch3']]
```

```
    df_list[0], a, my_figure0 = create_histo(name_list[0], df_list[0],
                                             my_figure0, col='blue')
```

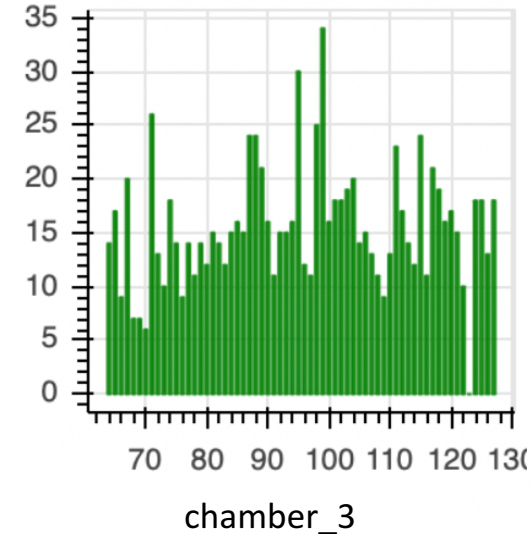


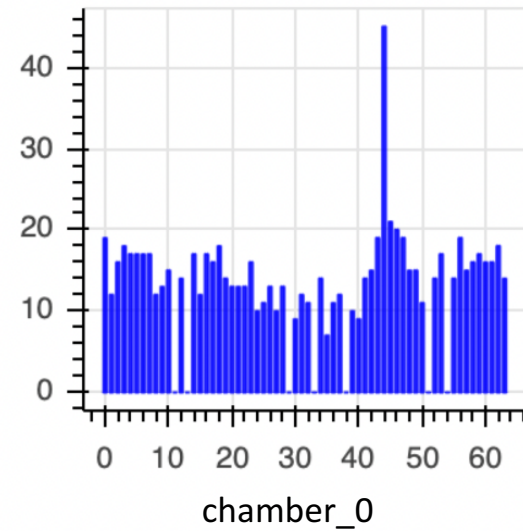
```
    df_list[1], b, my_figure1 = create_histo(name_list[1], df_list[1],
                                             my_figure1, col='red')
```

```
    df_list[2], c, my_figure2 = create_histo(name_list[2], df_list[2],
                                             my_figure2, col='yellow')
```






```
    df_list[3], d, my_figure3 = create_histo(name_list[3], df_list[3],
                                             my_figure3, col='green')
```



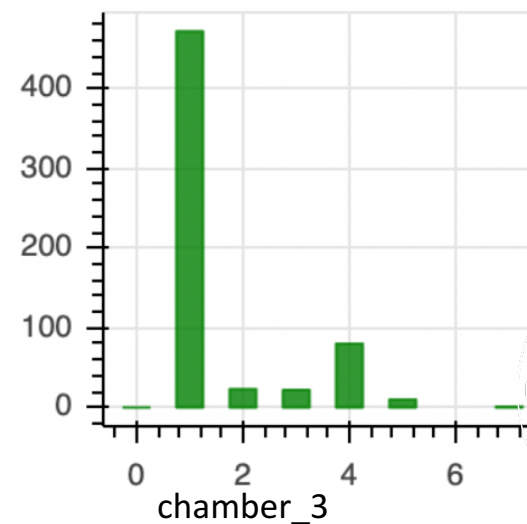

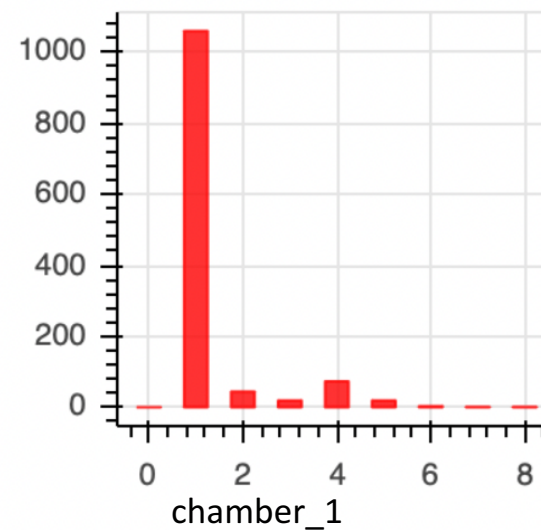

Plots of the counts of active TDC_CHANNEL



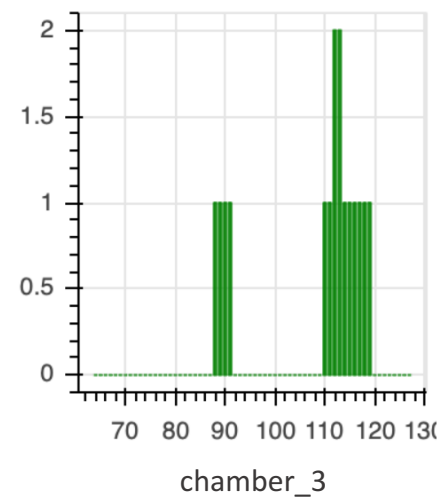



Plots of the total number of active TDC_CHANNEL in each ORBIT_CNT



Plots of the total number of active TDC_CHANNEL in each ORBIT_CNT



**Plots of the counts of active TDC_CHANNEL
only for those orbits with at least one scintillator signal in it**



To sum up:

- We set up a Spark and Kafka cluster with four workers and one broker on a multi-node VM cluster
- We implemented a streaming application using Kafka and Spark Structured Streaming
- We processed the data of the Drift Tubes detectors extracting the information of interest
- We were able to plot our results in a live updating dashboard using Bokeh



Thanks for the attention!

